

Scalable Matrix Architecture (Fundamentals and Vector-Matrix Facility)

José Moreira
IBM Corporation

Abel Bernabeu
Esperanto

Comments on the state-of-the-art for matrix computing

- Most high-performance matrix computations rely on the following kernel:

$$\mathbf{C}_{m \times n} = \mathbf{A}_{m \times K} \times \mathbf{B}_{K \times n}$$

$\mathbf{C}_{m \times n}$ is a register-resident matrix, with $m \approx n$, and
 $\mathbf{A}_{m \times K}$ and $\mathbf{B}_{K \times n}$ are streamed from memory, with $m \ll K$ and $n \ll K$

- Matrix $\mathbf{A}_{m \times K}$ is organized as a vector of K elements, each element a column vector of m elements of $\mathbf{A} = [\mathbf{A}^0 \quad \mathbf{A}^1 \quad \dots \quad \mathbf{A}^{K-1}]$


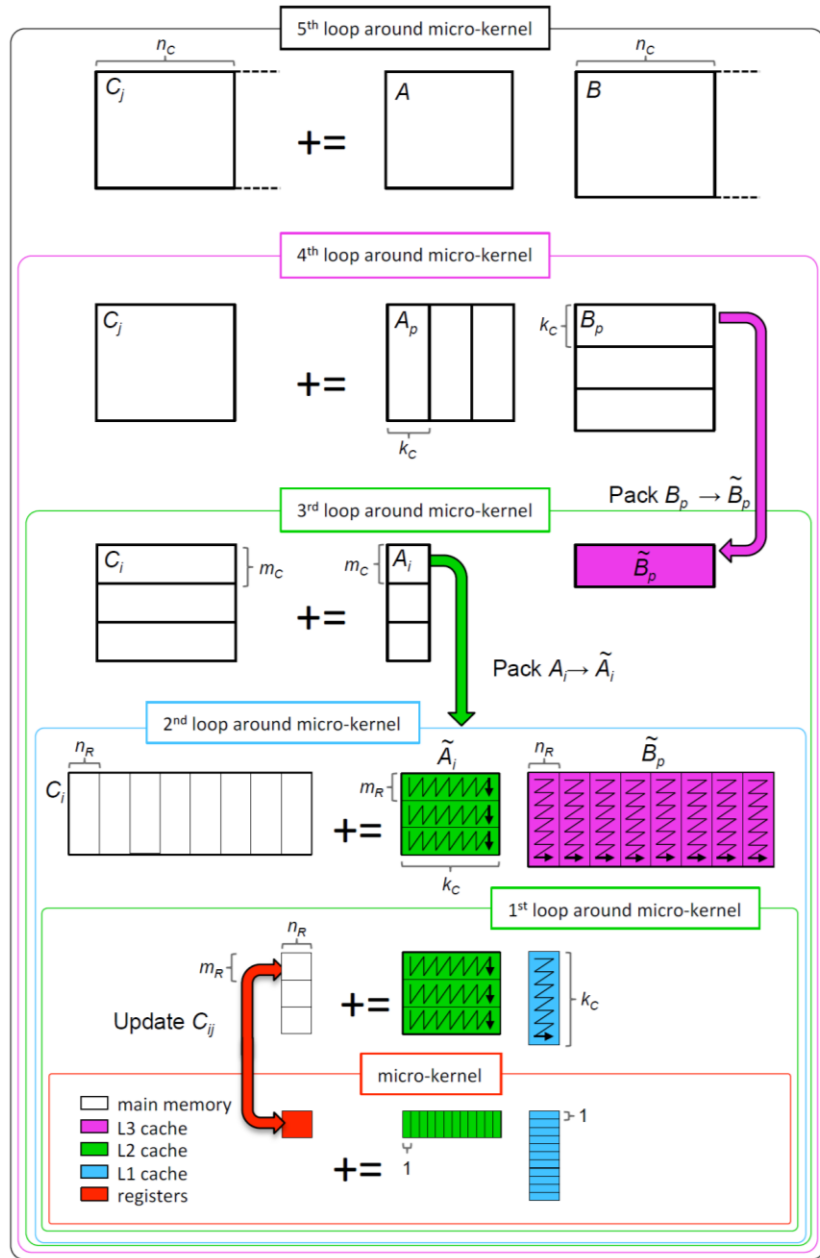
- Matrix $\mathbf{B}_{K \times n}$ is organized as a vector of K elements, each element a row of n

elements of $\mathbf{B} = \begin{bmatrix} \mathbf{B}_0 \\ \mathbf{B}_1 \\ \vdots \\ \mathbf{B}_{K-1} \end{bmatrix}$

- We will consider only the case of deterministic values of \mathbf{C} , as produced by the following algorithm:


$$\mathbf{C} \leftarrow 0; \text{for } (k = [0, K)) \mathbf{C} \leftarrow \mathbf{A}^k \times \mathbf{B}_k + \mathbf{C}$$

Anatomy of a high-performance matrix multiply



Performance bounds

- Let Δ be the latency (in cycles) of an elemental multiply-add operation
- Computing each element of \mathbf{C} , as described above, requires evaluating a dependence chain of K multiply-adds, and therefore takes time $K\Delta$

$$\mathbf{C} \leftarrow \mathbf{A}^k \times \mathbf{B}_k + \mathbf{C}$$


The diagram illustrates a feedback loop in the computation of \mathbf{C} . A blue bracket connects the \mathbf{C} on the right (the previous result) to the \mathbf{C} on the left (the current result being updated). Below the bracket is the symbol Δ , representing the latency of a single multiply-add operation.

- The computation of \mathbf{C} requires mnK multiply-adds
- The maximum computation rate that can be sustained is

$$R = \frac{mnK}{K\Delta} = mn/\Delta \quad \text{madds/cycle}$$

- This sets an upper bound on performance, dictated by the size of the \mathbf{C} panel that can be kept in registers (an architectural parameter) and the latency of a multiply-add (a design/technology parameter) – *Note*: integer arithmetic is more forgiving
- For modern server processor and 32-bit floating-point elements, a reasonable value of Δ is 4 cycles – the range $\Delta \in [2,8]$ should cover most cases

Some configurations (just to get a feel)

- Let the architected space consist of 32 vector registers of size L words (32-bit)
 - Then, $mn \leq 32L$ and the maximum computation rate is $R = 32L/4 = 8L$ madds/cycle (single-precision floating-point)
- The maximum computation rate of a vector-based matrix facility, using vector registers only, scales with L
- For $L = 4$, $R = 32$ madds/cycle
- For $L = 16$, $R = 128$ madds/cycle
- *Note:* In the above, we use all architected space for C , but in a load/store architecture we need to reserve space for A and B , so cut the bound in half
 - $mn \leq 16L$ and the maximum computation rate is $R = 16L/4 = 4L$ madds/cycle
 - For $L = 4$, $R = 16$ madds/cycle
 - For $L = 16$, $R = 64$ madds/cycle
- *Note:* $4L$ madds/cycle is the performance equivalent of 4 vector pipes operating concurrently

Conclusions

- The size of the architected register space imposes an upper bound on the performance of deterministic matrix multiply kernels
- If we use the architected scalable vector register space ($32L$ words) to hold the panel of C our performance upper bound, for a reasonable choice of latency parameters, is $8L$ madds/cycle – a more realistic value is $4L$ madds/cycle (single-precision floating-point)
- We advocate to pursue the vector register-based matrix extensions first, starting with three possible options:
 - Option A : 1 matrix/1 vector register $(R = 4L \text{ madds/cycle with } 4 \text{ matrix pipes})$
 - Option B : 1 matrix/ n vector registers $(R = 4L \text{ madds/cycle with } 1 \text{ matrix pipe})$
 - Option C : n matrices/1 vector register $(R = 4L \text{ madds/cycle with } 4/\lambda \text{ matrix pipes})$