

Matrix Tile Extension

Portable ISA for Vector-Integrated Matrix Units

Erich Focht (NEC), Marc Casas-Guix (BSC)

Scope and Principles

Provide a simple and portable way to use tensor units linked to vector units for performing the operations $\mathbf{C} = \mathbf{A} \cdot \mathbf{B}$ and $\mathbf{C} = \mathbf{C} + \mathbf{A} \cdot \mathbf{B}$ where \mathbf{A} , \mathbf{B} , \mathbf{C} are dense matrices of the dimensions, respectively, $m \times k$, $k \times n$, $m \times n$.


Portability

Few Commands

Simple Usage

The names of the instructions have been chosen to describe their function and are by no means “final”.

Accommodate Diversity



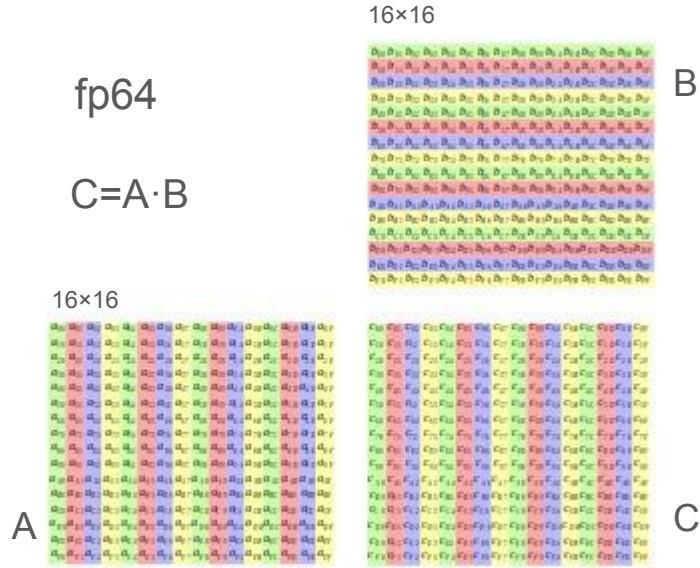
Short Vectors (e.g. VLEN=256)

- fp64: **2x2** or **4x1** matrix per VREG
- fp16: **4x4** or **8x2** matrix per VREG
- VRF: 4-8 read ports (?)
- 4 lanes (or less)

Long Vectors (e.g. VLEN=16384)

- fp64: **16x16**, **32x8**, **64x4**, **128x2**, **256x1** matrix
- fp16: **32x32**, **64x16**, **128x8**, **256x4**, **512x2**, **1024x1** matrix
- VRF: 5-16 read ports per lane, eg. 512 on SX-Aurora VE3
- 8 to 32 lanes
- many possible matrix multiplier geometries!

Square Tiles




Arithmetic Intensity:

Load: $2 \cdot 256 = 512$


Matmult: $2 \cdot 16 \cdot 16 \cdot 16 = 8192$ ops

Arithmetic intensity = $8192 / 512 = 16$

With 4x4 accumulators: a.i. = 64



Narrow Tiles




Arithmetic Intensity:

Load: $2 \cdot 4 \cdot 64 = 512$

Matmult: $2 \cdot 4 \cdot 4 \cdot 64 = 2048$ ops

Arithmetic intensity = $2048 / 512 = 4$

Narrow Tiles



Arithmetic Intensity:

Load: $64 \cdot 4 \cdot 2 = 512$

Matmult: $2 \cdot 64 \cdot 64 \cdot 4 = 32768$ ops


Arithmetic intensity = 64

A, B: 1 VREG each



C: 16 VREGs

Fundamentals

Matrix - Tile




Matrices



Tiles

Multiplication of the large matrices $C = A \cdot B$ is decomposed into sums of their multiplied tiles. Tiles are multiplied and added by the tensor unit.

Matrix - Tile




Maximum tile sizes, as supported by the tensor unit. Dimensions:

C-tiles: $m_{\max} \times n_{\max}$

A-tiles: $m_{\max} \times k_{\max}$

B-tiles: $k_{\max} \times n_{\max}$



Border tiles have smaller dimensions. Tensor units may or may not be able to operate on smaller tiles.

Tile - Matrix Unit - Vector Register

Matrix / tensor unit operates on tiles.



Tile dimensions are constrained by

- Tensor unit geometry: m_{\max} , k_{\max} , n_{\max}
 - May depend on SEW, data type
- Vector register size (VLEN, SEW)

Tiles are stored in vector registers.

- One or several (LMUL)
- Different tiles can have different LMUL

Matrix tile data layout in vector register is **not** specified and left to the implementation!




Tile Shape

A matrix tile is characterized by:

- VREG vector register number
- ROWS number of rows
- COLS number of columns
- SEW single element width
- MTYPE matrix type (A, B, C)
- (LMUL)
- (DTYPE data type (signed/unsigned/float/alt float))

***tile
shape***



MAXROWS, MAXCOLS depends on tensor unit geometry and tile type (A, B, C).

Tile Shape Data

The Tile Shape is a 32 bit entity that contains all information needed by the tensor unit to interpret correctly the tile data stored in a vector register or a vector register group.


- **rows/cols**: number of rows, columns of the tile
 - Values can vary widely: eg. 1 to 4096
 - Maximum VLEN=32768: 4096 elements in int8
 - Odd tail / border values should be handled as well
- **sew**: single element width
- **mtype**: tile type enum A, B, C
 - Load unit must recognize whether internal tile representation is in column-major-order or row-major-order.
- (**dtype**: Data type of tile: unsigned/signed/float/alt float)

Example:

```
typedef struct {  
    unsigned int rows : 12;  
    unsigned int cols : 12;  
    unsigned int sew : 3;  
    unsigned int mtype : 2;  
} TileShape_32b;
```

```
typedef enum {  
    MTYPE_A, MTYPE_B, MTYPE_C,  
} mtype_t;
```

Tensor Unit




Tensor Unit Configuration Needs

(1) Load matrix tile into VREG

Store matrix tile to memory

- Tile dimensions
- SEW
- Matrix type (A, B, C) for possible internal conversion

=> tile shape




Tensor Unit Configuration Needs

(1) Load matrix tile into VREG

Store matrix tile to memory

- Tile dimensions
- SEW
- Matrix type (A, B, C) for possible internal conversion

=> tile shape



(2) Multiply matrix tiles

- $m \times k \times n$,
- For each tile A, B, C:
 - SEW & DTYPE

or

- The 3 tile shapes of A, B, C
- (DTYPE, if not in tile shape)

(1) and (2) should be able to run simultaneously

Element-wise Matrix Operations

With matrix tiles stored in vector registers element-wise matrix operations can be done by using vector instructions.

- The effective or granted vector length depends on the size and shape of the tile.
- The tensor unit implementation can use masking to represent tiles smaller than the maximum tile size.
- Element-wise operations may need to use the corresponding, implementation-specific, vector mask.

Examples:

- Add two tiles of same shape
- Multiply each element of a tile with scalar value
- Initialize tile
- Apply activation function
- ...

Intrinsics/ISA

Core Logic

1. Get Tile Shape(s) for A, B, C

Iterate, strip-mine.

Like $vl = vsetvl(avl)$ but for 2D tiles

ask tensor unit for the
largest tile shape we need

tensor unit replies with
largest tile shape it supports

Core Logic

1. Get Tile Shape(s) for A, B, C

Iterate, strip-mine.

Like ***vl = vsetvl(avl)*** but for 2D tiles

ask tensor unit for the
largest tile shape we need

tensor unit replies with
largest tile shape it supports

2. Load A, B, C from memory

Like ***vleNN.v*** but for 2D tiles

Use tile shape data

Core Logic

1. Get Tile Shape(s) for A, B, C

Iterate, strip-mine.

Like ***vl = vsetvl(avl)*** but for 2D tiles

ask tensor unit for the
largest tile shape needed

tensor unit replies with
largest tile shape it supports

2. Load A, B, C from memory

Like ***vleNN.v*** but for 2D tiles

Use tile shape data

3. Multiply $C=C+A \cdot B$

Like ***vfmaccc.vv*** but for 2D tiles

Use tile shape data

Core Logic

1. Get Tile Shape(s) for A, B, C

Iterate, strip-mine.

Like ***vl = vsetvl(avl)*** but for 2D tiles

ask tensor unit for the
largest tile shape needed

tensor unit replies with
largest tile shape it supports

2. Load A, B, C from memory

Like ***vleNN.v*** but for 2D tiles

Use tile shape data

3. Multiply $C=C+A \cdot B$

Like ***vfmac.vv*** but for 2D tiles

Use tile shape data

4. Store C to memory

Like ***vseNN.v*** but for 2D tiles

Use tile shape data

Get Tile Shape

```
uint32_t tsget_T(int8_t mtype, uint32_t max_rows, uint32_t max_columns)
```

Retrieve the largest possible Tile Shape for a particular tile type A, B or C and the data type T. T is one of {i8, i16, i32, i64, u8, u16, u32, u64, **f16**, f32, f64}.

Equivalent to vsetvl intrinsic.

Arguments:

- **mtype**: matrix type enum value for A, B, C type.
- **max_rows**: maximum number of rows needed in the tile.
- **max_columns**: maximum number of columns needed in the tile.

Returns:

- A Tile Shape value that is supported by the tensor unit implementation.

Get Tile Shape

Remarks:

Only 1 VREG per tile case is considered.

The tile shape value removes the need to have state stored along the VREG containing a tile.

It eases optimization in the compiler like unrolling and reordering tile operations.

Requesting an unsupported data type T could either raise an exception or return a Zero tile shape value.


Letting the tensor implementation return the appropriate tile shape relieves us of the burden to manage various $m_{\max} \times k_{\max} \times n_{\max}$ combinations for different **data types** and **SEWs**.

Querying tile shapes simplifies iterating over the tiles when multiplying large matrices.

Get Tile Shape

Remarks:

Add LMUL to type T? Eg. *tsget_f16m4()* : Looks useful, needs further discussion.



Load Tile From Memory

```
vT_t tload_T(vT_t vd, uint64_t addr, uint32_t str_col, uint32_t str_row, uint32_t ts)
```

Load a tile shaped according to the tile shape ***ts*** from memory address ***addr*** into the vector register ***vd*** considering the column and row strides ***str_col***, ***str_row***.

T is the tile element data type (eg. f16 or f16m1).

Arguments:

- ***vd***: target vector register
- ***addr***: start address of the tile in memory
- ***str_col***: column stride of the tile data in memory
- ***str_row***: row stride of the tile data in memory
- ***ts***: tile shape formerly retrieved through call to ***tile_get_shape()***

Load Tile From Memory

```
vT_t tload_T(vT_t vd, uint64_t addr, uint32_t str_col, uint32_t str_row, uint32_t ts)
```

Load a tile shaped according to the tile shape ***ts*** from memory address ***addr*** into the vector register ***vd*** considering the column and row strides ***str_col***, ***str_row***.

T is the tile element data type (eg. f16 or f16m1).

Arguments:

- ***vd***: target vector register
- ***addr***: start address of the tile in memory
- ***str_col***: column stride of the tile data in memory
- ***str_row***: row stride of the tile data in memory
- ***ts***: tile shape formerly retrieved through call to ***tile_get_shape()***

Store Tile To Memory

```
tstore_T(vT_t vs, uint64_t addr, uint32_t str_col, uint32_t str_row, uint32_t ts)
```

Store a tile shaped according to the tile shape ***ts*** from the vector register ***vs*** to memory address ***addr*** considering the column and row strides ***str_col***, ***str_row***.

T is the tile element data type(eg. f32 or f32m1).

Arguments:

- ***vs***: source vector register
- ***addr***: start address of the tile in memory
- ***str_col***: column stride of the tile data
- ***str_row***: row stride of the tile data
- ***ts***: tile shape retrieved through ***tile_get_shape()***

Load / Store Tile

Remarks:

- The way how tile data is laid out in the vector register(s) is not specified and should remain implementation specific. It can be column-major-order or row-major-order or depend on the A, B, C matrix type.
- Passing both strides (row/column) helps the tensor unit transpose the tile when loading it, eg. if the internal representation requires column-major-order but data is row-major-order in memory. Order in memory doesn't matter.
- 2D slices of higher dimensional tensors can be loaded/stored.
- Strides in bytes or bytes in SEW?
- The intrinsics get transformed into (at least) 2 instructions: a ***vsetvli*** and the actual tile load instruction.

Multiply And Accumulate Tiles

$vT_{cm1_t} \text{ } tXmacc_T_c(vT_{cm1_t} \text{ } vc, vT_{AB}^{m1_t} \text{ } va, vT_{AB}^{m1_t} \text{ } vb, uint32_t \text{ } tc, uint32_t \text{ } ta, uint32_t \text{ } tb)$

Perform the multiply and accumulate operation $\mathbf{C}=\mathbf{C}+\mathbf{A}\cdot\mathbf{B}$ with tile \mathbf{A} in vector register \mathbf{va} , having shape \mathbf{ta} , \mathbf{B} stored in \mathbf{vb} , shape \mathbf{tb} and the in/out tile \mathbf{C} in \mathbf{vc} with shape \mathbf{tc} .

\mathbf{X} specifies operation type and widening. Eg. tfmacc, twmacc, tfwmacc, tqmacc, etc.

Arguments:

- vc : vector register that contains tile \mathbf{C} .
- va : vector register containing the tile \mathbf{A} .
- vb : vector register containing the tile \mathbf{B} .
- tc : tile shape of \mathbf{C} .
- ta : tile shape of \mathbf{A} .
- tb : tile shape of \mathbf{B} .

Multiply Tiles

Remarks:

- The intrinsics have 6 arguments: 3 x VREG + 3 x TileShape
 - Not feasible for assembler instructions, currently used encodings.
- Intrinsics get converted into at least two assembler instructions
 - Setup Tensor Unit: takes the three tile shapes as arguments
 - New instruction **tmulset**, similar to **vsetvli**.
 - Matrix Multiplication instruction: takes the three VREGs as arguments
- Handling LMUL is more complicated than in the vector case. LMUL in T_C with narrower T_A, T_B does not imply that $LMUL_{A,B} < LMUL_C$

Core Logic

1. Get Tile Shape(s) for A, B, C

Iterate, strip-mine.

Like ***vl = vsetvl(avl)*** but for 2D tiles

Intrinsics

tsget_f16m1

Assembler

(vsetvli)
tsget

2. Load A, B, C from memory

Like ***vleNN.v*** but for 2D tiles

tload_f16m1

vsetvli
tload

3. Multiply $C=C+A \cdot B$

Like ***vfmac.vv*** but for 2D tiles

tfmacc_f16m1

(vsetvli)
tmulset
tfmacc

4. Store C to memory

Like ***vseNN.v*** but for 2D tiles

tstore_f16m1

vsetvli
tstore

Tile Shape Helpers

Tile Shape Rows And Columns

```
uint16_t tsrows(uint32_t ts)
```

```
uint16_t tscols(uint32_t ts)
```

Extract **rows** and **columns** values from tile shape **ts**. This function can be a macro or compiler builtin, as it only does an AND and a SHIFT operation on the tile shape value. Expands to existing instructions.

Arguments:

- ts: tile shape.

Returns:

- Rows, columns values encoded in the tile shape variable.

Tile Shape Vector Length

```
size_t tsvl(uint32_t ts)
```

Return implementation specific vector length in elements that needs to be used for element-wise vector operations on the tile with the tile shape ***ts***.

Arguments:

- *ts*: tile shape

Returns:

- Vector length for element-wise operations on tile elements.

Tile Shape Vector Mask

```
vboolX_t tsmask(vboolX_t vd, uint32_t ts)
```

The implementation specific vector mask that needs to be used for element-wise vector operations on the tile with the tile shape **ts** is stored in the destination vector register **vd**. Typically masks are needed when the tile shape doesn't have the maximum size.

Arguments:

- **vd**: destination vector register for vector mask.
- **ts**: tile shape.


Returns:

- Vector mask for element-wise operations on tile elements.


Example:

4 x 64


FP64



4 x 63



3 x 64



Assembler

tsget	Tile shape get	
tload	Tile load	
tstore	Tile store	
tmac	Tile multiply-accumulate	tfmacc, tfwmacc, twmacc, tqmacc, tfqmacc, tmacu,...
tmulset	Tile multiplier set	
tsvl	Tile shape vector length	
tmask	Tile shape mask	

NOTE: Adding data type information to Tile Shape reduces number of *tmac* instruction variants.

Tensor Unit CSR (example)

```
struct TensorUnitCSR {  
    // - tile shapes -  
    unsigned int m : 12;  
    unsigned int k : 12;  
    unsigned int n : 12;  
    unsigned int sew_a : 3;  
    unsigned int sew_b : 3;  
    unsigned int sew_c : 3;  
    // - data type -  
    unsigned int float : 1;  
    unsigned int alt_float : 1;  
    unsigned int unsig : 1;  
    // - further config -  
    unsigned int sparse : 1;  
    // - status -  
    unsigned int status : X;  
    // - implementation specific -  
    unsigned int impl : Y;  
};
```

Examples

GEMM $C = A \cdot B$

Matrix data types: C: FP32 A,B: FP32

```
01: for m = 0, M, m += tm do
02:     for n = 0, N, n += tn do
03:         tc = tsget_f32(MTYPE_C, M - m, N - n)
04:         tm = tsrows(tc)
05:         tn = tscols(tc)
06:         gvl = tsvl(tc)
07:         c = __riscv_vfmv_v_f_f32m1(0.0f, gvl)
08:         for k = 0, K, k += tscols(ta) do
09:             ta = tsget_f32(MTYPE_A, tm, K - k)
10:             tb = tsget_f32(MTYPE_B, K - k, tn)
11:             a = tload_f32(&A[m * LDA + k], LDA, 1, ta)
12:             b = tload_f32(&B[k * LDB + n], LDB, 1, tb)
13:             c = tfmacc_f32(c, a, b, tc, ta, tb)
14:         tstore_f32(c, &C[m * LDC + n], LDC, 1, tc)
```

GEMM $C = A \cdot B$ – Using $\text{LMUL} > 1$

Matrix data types: C: FP32 A,B: FP32

```
01: for m = 0, M, m += tm do
02:   for n = 0, N, n += tn do
03:     tc = tsget_f32m4(MTYPE_C, M - m, N - n)
04:     tm = tsrows(tc)
05:     tn = tscols(tc)
06:     gvl = tsvl(tc)
07:     c = __riscv_vfmv_v_f_f32m4(0.0f, gvl)
08:     for k = 0, K, k += tscols(ta) do
09:       ta = tsget_f32m1(MTYPE_A, tm, K - k)
10:       tb = tsget_f32m4(MTYPE_B, K - k, tn)
11:       a = tload_f32m1(&A[m * LDA + k], LDA, 1, ta)
12:       b = tload_f32m4(&B[k * LDB + n], LDB, 1, tb)
13:       c = tfmacc_f32(c, a, b, tc, ta, tb)
14:     tstore_f32m4(c, &C[m * LDC + n], LDC, 1, tc)
```



GEMM $C = \alpha \cdot A \cdot B + \beta \cdot C$

Matrix data types: C: FP32 A,B: FP16

```
01: for m = 0, M, m += tm do
02:     for n = 0, N, n += tn do
03:         tc = tsget_f32(MTYPE_C, M - m, N - n)
04:         tm = tsrows(tc)
05:         tn = tscols(tc)
06:         gvl = tsvl(tc)
07:         cmask = tsmask(tc)
09:         c = __riscv_vfmv_v_f_f32m1(0.0f, gvl)
10:         for k = 0, K, k += tscols(ta) do
11:             ta = tsget_f16(MTYPE_A, tm, K - k)
12:             tb = tsget_f16(MTYPE_B, K - k, tn)
13:             a = tload_f16(&A[m * LDA + k], LDA, 1, ta)
14:             b = tload_f16(&B[k * LDB + n], LDB, 1, tb)
15:             c = tfwmacc_f32(c, a, b, tc, ta, tb)
16:             t = tload_f32(&C[m * LDC + n], LDC, 1, tc)
17:             c = __riscv_vfmul_vf_f32m1(c, alpha, gvl, cmask)
18:             c = __riscv_vfadd_vf_f32m1(t, beta, gvl, cmask)
19:             tile_store_fp32(c, &C[m * LDC + n], LDC, 1, tc)
```


TODO

TODO

- Discover supported data formats (!)
- Multiple VREG Tiles
 - Code for square matrices pattern (A) with 4x4 C-type + 4 A-type + 4 B-type = 24 VREGs
 - Not optimal for narrow matrices (B): 16 C-type VREGs + 1 A-type + 1 B-type = 18 VREGs
 - Pattern (A) probably still good enough to reach peak performance/throughput!
 - Add new Multi-Tile-Shape to describe optimal topology?
 - Eg. add rows, columns of tiles for optimal performance
 - Generate code dynamically at runtime for presented topology?
 - Or let compiler generate alternative code for extreme cases and chose path according to retrieved shape?
 - LMUL in tile shape: not sufficient for 2D tile patterns
 - If LMUL is hidden in an argument (ts): not possible to do proper register allocation.
 - Helps reduce number of instructions.
- Sparse matrix operations
- Other matrix operations

THE END