

Integrated Matrix Extension ISA

Proposal G

Erich Focht (Openchip)

Building on «Proposal A» with Marc Casas (BSC)

With contributions from Philipp Tomsich (VRULL)



24/02/2025

MATRIX IN MEMORY


$X_0^0 X_0^1 X_0^2 X_0^3 X_0^4 X_0^5 X_0^6 X_0^7 X_0^8 X_0^9 X_0^{10} X_0^{11} X_0^{12} X_0^{13} X_0^{14} X_0^{15} X_0^{16} X_0^{17} X_0^{18} X_0^{19} X_0^{20} X_0^{21} X_0^{22} X_0^{23} X_0^{24} X_0^{25} X_0^{26} X_0^{27} X_0^{28} X_0^{29} X_0^{30} X_0^{31}$
 $X_1^0 X_1^1 X_1^2 X_1^3 X_1^4 X_1^5 X_1^6 X_1^7 X_1^8 X_1^9 X_1^{10} X_1^{11} X_1^{12} X_1^{13} X_1^{14} X_1^{15} X_1^{16} X_1^{17} X_1^{18} X_1^{19} X_1^{20} X_1^{21} X_1^{22} X_1^{23} X_1^{24} X_1^{25} X_1^{26} X_1^{27} X_1^{28} X_1^{29} X_1^{30} X_1^{31}$
 $X_2^0 X_2^1 X_2^2 X_2^3 X_2^4 X_2^5 X_2^6 X_2^7 X_2^8 X_2^9 X_2^{10} X_2^{11} X_2^{12} X_2^{13} X_2^{14} X_2^{15} X_2^{16} X_2^{17} X_2^{18} X_2^{19} X_2^{20} X_2^{21} X_2^{22} X_2^{23} X_2^{24} X_2^{25} X_2^{26} X_2^{27} X_2^{28} X_2^{29} X_2^{30} X_2^{31}$
 $X_3^0 X_3^1 X_3^2 X_3^3 X_3^4 X_3^5 X_3^6 X_3^7 X_3^8 X_3^9 X_3^{10} X_3^{11} X_3^{12} X_3^{13} X_3^{14} X_3^{15} X_3^{16} X_3^{17} X_3^{18} X_3^{19} X_3^{20} X_3^{21} X_3^{22} X_3^{23} X_3^{24} X_3^{25} X_3^{26} X_3^{27} X_3^{28} X_3^{29} X_3^{30} X_3^{31}$
 $X_4^0 X_4^1 X_4^2 X_4^3 X_4^4 X_4^5 X_4^6 X_4^7 X_4^8 X_4^9 X_4^{10} X_4^{11} X_4^{12} X_4^{13} X_4^{14} X_4^{15} X_4^{16} X_4^{17} X_4^{18} X_4^{19} X_4^{20} X_4^{21} X_4^{22} X_4^{23} X_4^{24} X_4^{25} X_4^{26} X_4^{27} X_4^{28} X_4^{29} X_4^{30} X_4^{31}$
 $X_5^0 X_5^1 X_5^2 X_5^3 X_5^4 X_5^5 X_5^6 X_5^7 X_5^8 X_5^9 X_5^{10} X_5^{11} X_5^{12} X_5^{13} X_5^{14} X_5^{15} X_5^{16} X_5^{17} X_5^{18} X_5^{19} X_5^{20} X_5^{21} X_5^{22} X_5^{23} X_5^{24} X_5^{25} X_5^{26} X_5^{27} X_5^{28} X_5^{29} X_5^{30} X_5^{31}$
 $X_6^0 X_6^1 X_6^2 X_6^3 X_6^4 X_6^5 X_6^6 X_6^7 X_6^8 X_6^9 X_6^{10} X_6^{11} X_6^{12} X_6^{13} X_6^{14} X_6^{15} X_6^{16} X_6^{17} X_6^{18} X_6^{19} X_6^{20} X_6^{21} X_6^{22} X_6^{23} X_6^{24} X_6^{25} X_6^{26} X_6^{27} X_6^{28} X_6^{29} X_6^{30} X_6^{31}$
 $X_7^0 X_7^1 X_7^2 X_7^3 X_7^4 X_7^5 X_7^6 X_7^7 X_7^8 X_7^9 X_7^{10} X_7^{11} X_7^{12} X_7^{13} X_7^{14} X_7^{15} X_7^{16} X_7^{17} X_7^{18} X_7^{19} X_7^{20} X_7^{21} X_7^{22} X_7^{23} X_7^{24} X_7^{25} X_7^{26} X_7^{27} X_7^{28} X_7^{29} X_7^{30} X_7^{31}$
 $X_8^0 X_8^1 X_8^2 X_8^3 X_8^4 X_8^5 X_8^6 X_8^7 X_8^8 X_8^9 X_8^{10} X_8^{11} X_8^{12} X_8^{13} X_8^{14} X_8^{15} X_8^{16} X_8^{17} X_8^{18} X_8^{19} X_8^{20} X_8^{21} X_8^{22} X_8^{23} X_8^{24} X_8^{25} X_8^{26} X_8^{27} X_8^{28} X_8^{29} X_8^{30} X_8^{31}$
 $X_9^0 X_9^1 X_9^2 X_9^3 X_9^4 X_9^5 X_9^6 X_9^7 X_9^8 X_9^9 X_9^{10} X_9^{11} X_9^{12} X_9^{13} X_9^{14} X_9^{15} X_9^{16} X_9^{17} X_9^{18} X_9^{19} X_9^{20} X_9^{21} X_9^{22} X_9^{23} X_9^{24} X_9^{25} X_9^{26} X_9^{27} X_9^{28} X_9^{29} X_9^{30} X_9^{31}$
 $X_{10}^0 X_{10}^1 X_{10}^2 X_{10}^3 X_{10}^4 X_{10}^5 X_{10}^6 X_{10}^7 X_{10}^8 X_{10}^9 X_{10}^{10} X_{10}^{11} X_{10}^{12} X_{10}^{13} X_{10}^{14} X_{10}^{15} X_{10}^{16} X_{10}^{17} X_{10}^{18} X_{10}^{19} X_{10}^{20} X_{10}^{21} X_{10}^{22} X_{10}^{23} X_{10}^{24} X_{10}^{25} X_{10}^{26} X_{10}^{27} X_{10}^{28} X_{10}^{29} X_{10}^{30} X_{10}^{31}$
 $X_{11}^0 X_{11}^1 X_{11}^2 X_{11}^3 X_{11}^4 X_{11}^5 X_{11}^6 X_{11}^7 X_{11}^8 X_{11}^9 X_{11}^{10} X_{11}^{11} X_{11}^{12} X_{11}^{13} X_{11}^{14} X_{11}^{15} X_{11}^{16} X_{11}^{17} X_{11}^{18} X_{11}^{19} X_{11}^{20} X_{11}^{21} X_{11}^{22} X_{11}^{23} X_{11}^{24} X_{11}^{25} X_{11}^{26} X_{11}^{27} X_{11}^{28} X_{11}^{29} X_{11}^{30} X_{11}^{31}$
 $X_{12}^0 X_{12}^1 X_{12}^2 X_{12}^3 X_{12}^4 X_{12}^5 X_{12}^6 X_{12}^7 X_{12}^8 X_{12}^9 X_{12}^{10} X_{12}^{11} X_{12}^{12} X_{12}^{13} X_{12}^{14} X_{12}^{15} X_{12}^{16} X_{12}^{17} X_{12}^{18} X_{12}^{19} X_{12}^{20} X_{12}^{21} X_{12}^{22} X_{12}^{23} X_{12}^{24} X_{12}^{25} X_{12}^{26} X_{12}^{27} X_{12}^{28} X_{12}^{29} X_{12}^{30} X_{12}^{31}$
 $X_{13}^0 X_{13}^1 X_{13}^2 X_{13}^3 X_{13}^4 X_{13}^5 X_{13}^6 X_{13}^7 X_{13}^8 X_{13}^9 X_{13}^{10} X_{13}^{11} X_{13}^{12} X_{13}^{13} X_{13}^{14} X_{13}^{15} X_{13}^{16} X_{13}^{17} X_{13}^{18} X_{13}^{19} X_{13}^{20} X_{13}^{21} X_{13}^{22} X_{13}^{23} X_{13}^{24} X_{13}^{25} X_{13}^{26} X_{13}^{27} X_{13}^{28} X_{13}^{29} X_{13}^{30} X_{13}^{31}$
 $X_{14}^0 X_{14}^1 X_{14}^2 X_{14}^3 X_{14}^4 X_{14}^5 X_{14}^6 X_{14}^7 X_{14}^8 X_{14}^9 X_{14}^{10} X_{14}^{11} X_{14}^{12} X_{14}^{13} X_{14}^{14} X_{14}^{15} X_{14}^{16} X_{14}^{17} X_{14}^{18} X_{14}^{19} X_{14}^{20} X_{14}^{21} X_{14}^{22} X_{14}^{23} X_{14}^{24} X_{14}^{25} X_{14}^{26} X_{14}^{27} X_{14}^{28} X_{14}^{29} X_{14}^{30} X_{14}^{31}$
 $X_{15}^0 X_{15}^1 X_{15}^2 X_{15}^3 X_{15}^4 X_{15}^5 X_{15}^6 X_{15}^7 X_{15}^8 X_{15}^9 X_{15}^{10} X_{15}^{11} X_{15}^{12} X_{15}^{13} X_{15}^{14} X_{15}^{15} X_{15}^{16} X_{15}^{17} X_{15}^{18} X_{15}^{19} X_{15}^{20} X_{15}^{21} X_{15}^{22} X_{15}^{23} X_{15}^{24} X_{15}^{25} X_{15}^{26} X_{15}^{27} X_{15}^{28} X_{15}^{29} X_{15}^{30} X_{15}^{31}$

MATRIX IN MEMORY

stride 1 access
direction



row-major order




column-major order

MEMORY TO VREG TILE CONTENT


Matrix in memory – FP32

VLEN=256



MEMORY TO VREG TILE CONTENT

Matrix in memory – FP32



VLEN=256

X_0^{16}	X_0^{17}	X_0^{18}	X_0^{19}	X_0^{20}	X_0^{21}	X_0^{22}	X_0^{23}
------------	------------	------------	------------	------------	------------	------------	------------


X_4^{16}	X_4^{17}	X_4^{18}	X_4^{19}
X_5^{16}	X_5^{17}	X_5^{18}	X_5^{19}

Possible layout in VREG
(specific to implementation)

X_4^{16}	X_4^{17}	X_4^{18}	X_4^{19}	X_5^{16}	X_5^{17}	X_5^{18}	X_5^{19}
------------	------------	------------	------------	------------	------------	------------	------------

MEMORY TO VREG TILE CONTENT

Matrix in memory – FP32



VLEN=256

X_0^{16}	X_0^{17}	X_0^{18}	X_0^{19}	X_0^{20}	X_0^{21}	X_0^{22}	X_0^{23}
------------	------------	------------	------------	------------	------------	------------	------------

X_4^{16}	X_4^{17}	X_4^{18}	X_4^{19}
X_5^{16}	X_5^{17}	X_5^{18}	X_5^{19}

X_8^{16}	X_8^{17}	X_8^{18}	X_8^{19}	X_8^{20}	X_8^{21}	X_8^{22}	X_8^{23}
X_9^{16}	X_9^{17}	X_9^{18}	X_9^{19}	X_9^{20}	X_9^{21}	X_9^{22}	X_9^{23}
X_{10}^{16}	X_{10}^{17}	X_{10}^{18}	X_{10}^{19}	X_{10}^{20}	X_{10}^{21}	X_{10}^{22}	X_{10}^{23}
X_{11}^{16}	X_{11}^{17}	X_{11}^{18}	X_{11}^{19}	X_{11}^{20}	X_{11}^{21}	X_{11}^{22}	X_{11}^{23}

X_8^{16}	X_8^{17}	X_8^{18}	X_8^{19}	X_8^{20}	X_8^{21}	X_8^{22}	X_8^{23}
X_9^{16}	X_9^{17}	X_9^{18}	X_9^{19}	X_9^{20}	X_9^{21}	X_9^{22}	X_9^{23}
X_{10}^{16}	X_{10}^{17}	X_{10}^{18}	X_{10}^{19}	X_{10}^{20}	X_{10}^{21}	X_{10}^{22}	X_{10}^{23}
X_{11}^{16}	X_{11}^{17}	X_{11}^{18}	X_{11}^{19}	X_{11}^{20}	X_{11}^{21}	X_{11}^{22}	X_{11}^{23}

X_8^{16}	X_8^{17}	X_8^{18}	X_8^{19}	X_8^{20}	X_8^{21}	X_8^{22}	X_8^{23}
X_9^{16}	X_9^{17}	X_9^{18}	X_9^{19}	X_9^{20}	X_9^{21}	X_9^{22}	X_9^{23}
X_{10}^{16}	X_{10}^{17}	X_{10}^{18}	X_{10}^{19}	X_{10}^{20}	X_{10}^{21}	X_{10}^{22}	X_{10}^{23}
X_{11}^{16}	X_{11}^{17}	X_{11}^{18}	X_{11}^{19}	X_{11}^{20}	X_{11}^{21}	X_{11}^{22}	X_{11}^{23}


Possible layout in VREG
(specific to implementation)

X_4^{16}	X_4^{17}	X_4^{18}	X_4^{19}	X_5^{16}	X_5^{17}	X_5^{18}	X_5^{19}
------------	------------	------------	------------	------------	------------	------------	------------

Multiple
VREGs

MEMORY TO VREG TILE CONTENT

Matrix in memory – FP32



VLEN=256

$$\begin{bmatrix} X_0^{16} & X_0^{17} & X_0^{18} & X_0^{19} & X_0^{20} & X_0^{21} & X_0^{22} & X_0^{23} \end{bmatrix}$$

RMUL=1
CMUL=1

Possible layout in VREG
(specific to implementation)

$$\begin{bmatrix} X_4^{16} & X_4^{17} & X_4^{18} & X_4^{19} \\ X_5^{16} & X_5^{17} & X_5^{18} & X_5^{19} \end{bmatrix}$$

$$\begin{bmatrix} X_4^{16} & X_4^{17} & X_4^{18} & X_4^{19} & X_5^{16} & X_5^{17} & X_5^{18} & X_5^{19} \end{bmatrix}$$

$$\begin{bmatrix} X_8^{16} & X_8^{17} & X_8^{18} & X_8^{19} & X_8^{20} & X_8^{21} & X_8^{22} & X_8^{23} \\ X_9^{16} & X_9^{17} & X_9^{18} & X_9^{19} & X_9^{20} & X_9^{21} & X_9^{22} & X_9^{23} \\ X_{10}^{16} & X_{10}^{17} & X_{10}^{18} & X_{10}^{19} & X_{10}^{20} & X_{10}^{21} & X_{10}^{22} & X_{10}^{23} \\ X_{11}^{16} & X_{11}^{17} & X_{11}^{18} & X_{11}^{19} & X_{11}^{20} & X_{11}^{21} & X_{11}^{22} & X_{11}^{23} \end{bmatrix}$$

RMUL=4
CMUL=1

$$\begin{bmatrix} X_8^{16} & X_8^{17} & X_8^{18} & X_8^{19} & X_8^{20} & X_8^{21} & X_8^{22} & X_8^{23} \\ X_9^{16} & X_9^{17} & X_9^{18} & X_9^{19} & X_9^{20} & X_9^{21} & X_9^{22} & X_9^{23} \\ X_{10}^{16} & X_{10}^{17} & X_{10}^{18} & X_{10}^{19} & X_{10}^{20} & X_{10}^{21} & X_{10}^{22} & X_{10}^{23} \\ X_{11}^{16} & X_{11}^{17} & X_{11}^{18} & X_{11}^{19} & X_{11}^{20} & X_{11}^{21} & X_{11}^{22} & X_{11}^{23} \end{bmatrix}$$

RMUL=2
CMUL=2

RMUL=1
CMUL=4

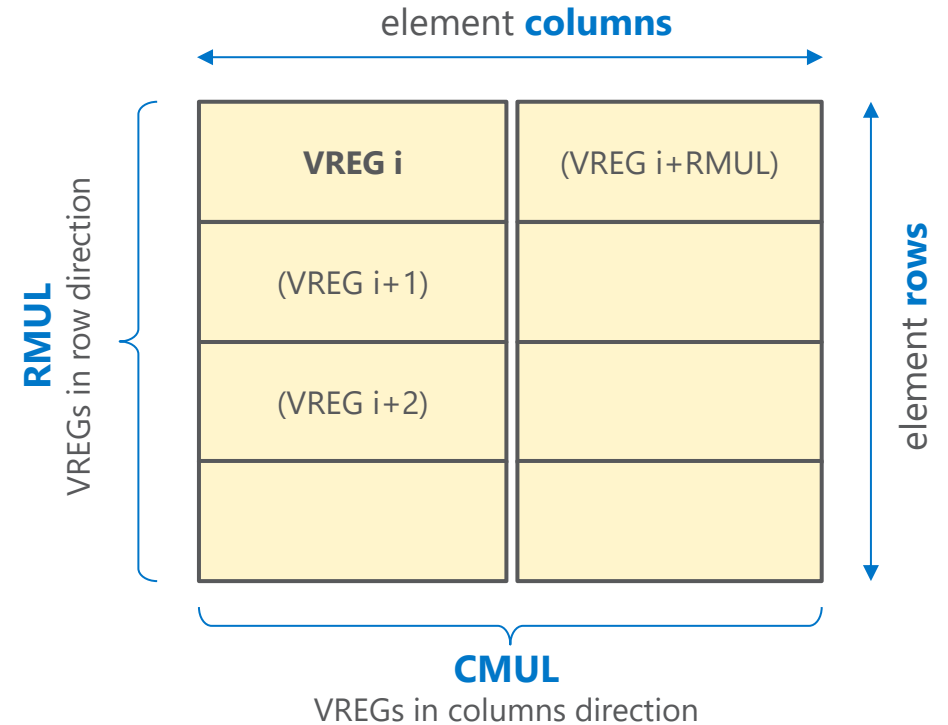
$$\begin{bmatrix} X_8^{16} & X_8^{17} & X_8^{18} & X_8^{19} & X_8^{20} & X_8^{21} & X_8^{22} & X_8^{23} \\ X_9^{16} & X_9^{17} & X_9^{18} & X_9^{19} & X_9^{20} & X_9^{21} & X_9^{22} & X_9^{23} \\ X_{10}^{16} & X_{10}^{17} & X_{10}^{18} & X_{10}^{19} & X_{10}^{20} & X_{10}^{21} & X_{10}^{22} & X_{10}^{23} \\ X_{11}^{16} & X_{11}^{17} & X_{11}^{18} & X_{11}^{19} & X_{11}^{20} & X_{11}^{21} & X_{11}^{22} & X_{11}^{23} \end{bmatrix}$$

Vector Register Matrix Tiles

POSSIBLE CONVENTION FOR VREG TILES

A matrix tile is stored in a vector register or a vector register group.

- RMUL (row multiplier) stacks VREGs vertically. Power of 2. Maximum 16.
- CMUL (column multiplier) multiplies number of VREG columns. Power of 2. Maximum 16.
- Specify only **VREG *i*** as argument to matrix unit commands.
- VREG group content metadata: „tile shape“



```
// example
struct TileShape {
    unsigned int rows : MM;
    unsigned int cols : NN;
    unsigned int rmul  : 3;
    unsigned int cmul  : 3;
    unsigned int sew   : 3;
    unsigned int mtype : 2;
    unsigned int cmo   : 1;
};
```


INSTRUCTIONS OVERVIEW

Logical instructions

GET Tile Shape

vmts rd, rs1, sew, mtype, order

LOAD Tile

vm1t vd, (rs2), rs1, ts

MULTIPLY Tiles

vmmacc.vv v_C, v_A, v_B, ts_C, ts_A, ts_B

STORE Tile

vmst vd, (rs2), rs1, ts

INSTRUCTIONS OVERVIEW

vmts rd, rs1, sew, mtype, order

GET supported tile shape

rd→tile shape, rs1→requested rows, cols

vmmset tc, ta, tb

SET/configure matrix multiplier

ta, tb, tc : tile shapes of A, B, C

vm1t vd, (rs2), rs1

LOAD matrix tile from memory

vd→VREG grp, rs2→base address, rs1→tile type, strides,...

vmst vd, (rs2), rs1

STORE matrix tile to memory

vd→VREG grp, rs2→base address, rs1→tile type, strides,...

vmmacc.vv vd, vs1, vs2

matrix **MULTIPLY** and accumulate vd→C, vs1→A, vs2→B

THE FLOW



QUERY

SW can request impl. limits

- max RMUL, CMUL for given SEW, matrix type, ordering, widening
- Instruction or runtime mechanism
- No effect on HW, read-only
- Assume that for RMUL, CMUL all lower values are supported, too.
4 → 2, 1

EXECUTION PLAN

SW determines appropriate settings for execution strategy

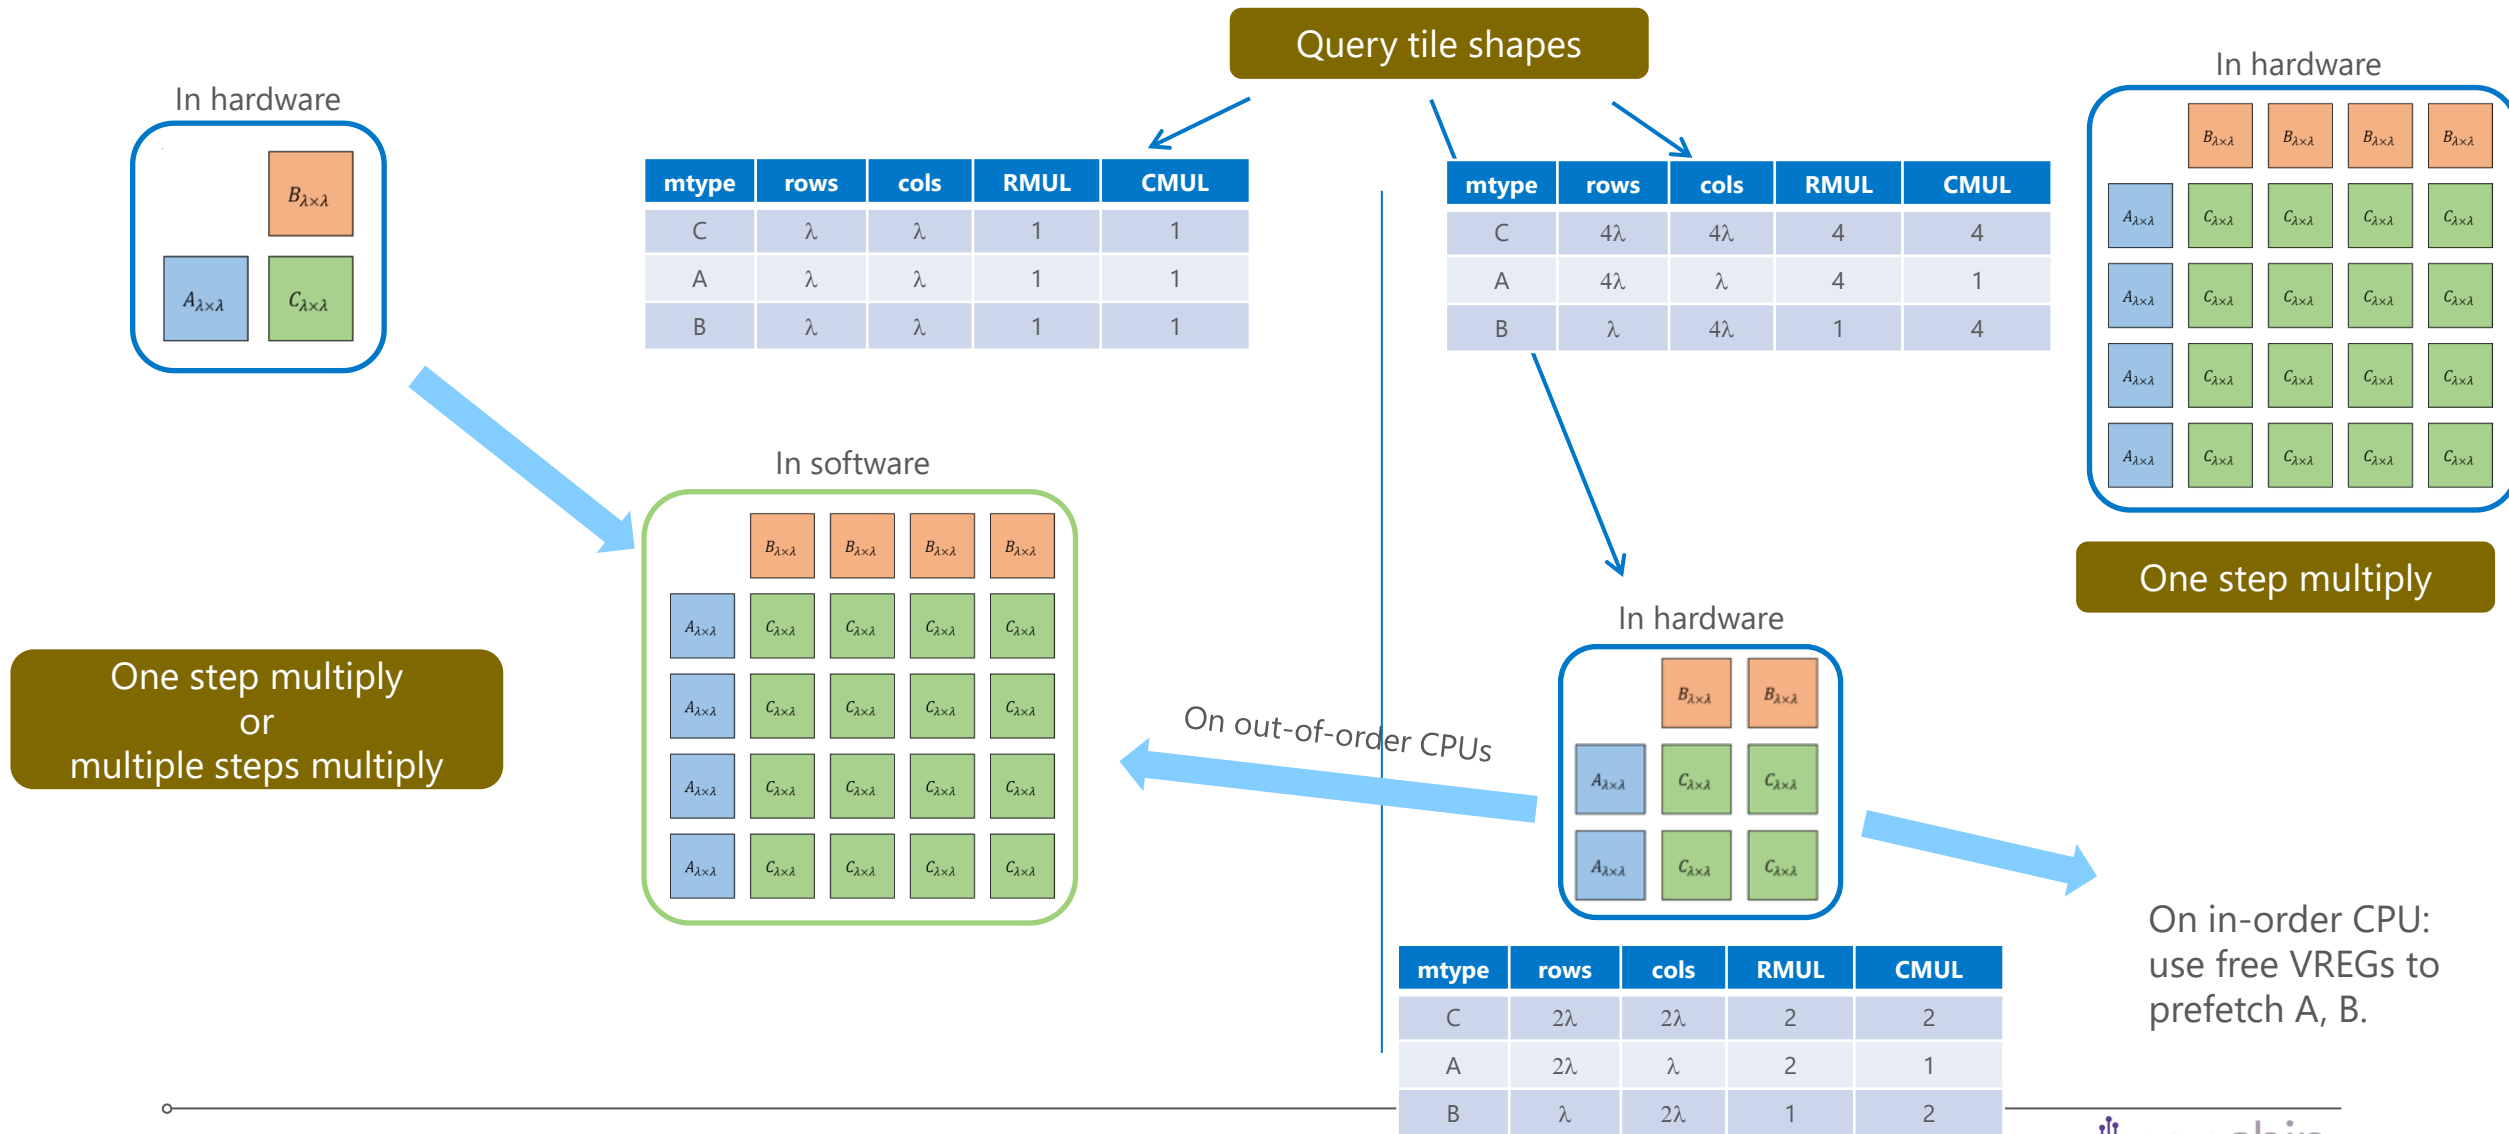
- Register allocation
- Memory hierarchy
- Prefetching or eager loading ahead of use
- Can be IF blocks selecting among multiple optimized implementations.

CONFIGURE


Write values corresponding to the execution plan
e.g. into CSRs, configuring the matrix unit

EXECUTE

MAPPING TO IME OPTION A




MAPPING TO IME OPTION C




mtype	rows	cols	RMUL	CMUL
C	16	16	8	2
A	16	2	2	1
B	2	16	1	2

One step multiply



mtype	rows	cols	RMUL	CMUL
C	8	8	4	1
A	8	2	1	1
B	2	8	1	1

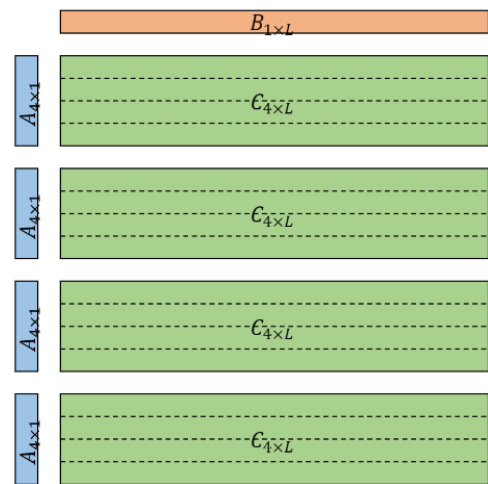


mtype	rows	cols	RMUL	CMUL
C	8	8	1	1
A	2	2	1	1
B	2	8	1	1

Using only a segment of A!

One step multiply
or
4 steps multiply


MAPPING TO IME OPTION B




mtype	rows	cols	RMUL	CMUL
C	16	L	16	1
A	16	1	1	1
B	1	L	1	1

1 multiply

mtype	rows	cols	RMUL	CMUL
C	4	L	4	1
A	4	1	1	1
B	1	L	1	1




... 4 multiplies




Stacking several A and B registers seems more efficient.

MAPPING TO IME OPTION D

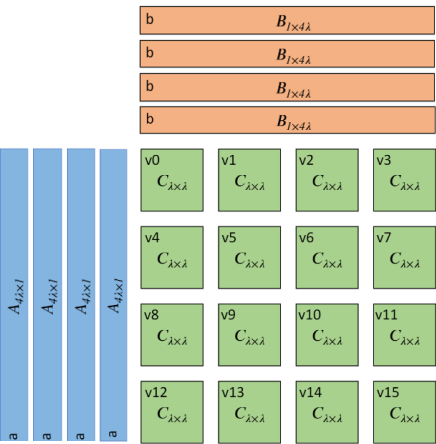


mtype	rows	cols	RMUL	CMUL
C	4λ	4λ	4	4
A	4λ	1	1	1
B	1	4λ	1	1



a bit wasteful...


a bit better?



mtype	rows	cols	RMUL	CMUL
C	4λ	4λ	4	4
A	4λ	1	1	4
B	1	4λ	4	1

This implementation option is excellent for column-major-order A and row-major-order B (not what the evaluation was asking for...)


MAPPING IME OPTION E



mtype	rows	cols	RMUL	CMUL
C	8	8	4	1
A	8	32	1	4
B	32	8	4	1


mtype	rows	cols	RMUL	CMUL
C	16	16	8	2
A	16	32	2	4
B	32	16	4	2

MAPPING IME OPTION F



mtype	rows	cols	RMUL	CMUL	order
C	1	8	1	1	rmo
A	1	VLEN/SEW	1	1	rmo
B	VLEN/SEW	8	8	1	cmo


Straight forward ...



*Segment handling needs vslideup or similar ...
But is it really needed if B is weight stationary?*

NOTE: a **weight stationary** implementation needs a different order of strip-mining loops! N, K, M

IME TG - SEMIDYNAMICS




mtype	rows	cols	RMUL	CMUL
C	16	64	16	1
A	16	16	2	1
B	16	64	8	1

Straight forward ...

COLUMN-MAJOR-ORDER A MATRIX – HPC DESIGN POINT

FP16, INT16
widen 1x



mtype	rows	cols	RMUL	CMUL
C	128	128	16	1
A	128	8	1	1
B	8	128	1	1

QUERY MATRIX TILE SHAPE

New instruction: Vector Float Matrix Tile Shape:

vmts rd, rs1, sew, mtype, order

Retrieve the Tile Shape supported by the matrix extension for a particular tile type A, B or C which matches the desired maximum rows/columns. Equivalent to *vsetvl*.

Arguments:

- **rd**: destination scalar register, contains **tile shape** result
- **rs1**: source register contains desired (maximum) **rows, columns** shape in memory (16 bit each, for example)
- **sew**: desired single element width, like *vsetvl*: *e8, e16, e32, e64*
- **mtype**: matrix type for A, B, C type: *mta, mtb, mtc*
- **order**: order in memory, column- or row-major-order: *cmo, rmo*

Returns:

- A **TileShape** value that is supported by the matrix unit implementation.

vmts
vfmts
vfwmts
vfqmts

```
// example
struct TileShape {
    unsigned int rows : MM;
    unsigned int cols : NN;
    unsigned int rmul  : 3;
    unsigned int cmul  : 3;
    unsigned int sew   : 3;
    unsigned int mtype : 2;
    unsigned int cmo   : 1;
};
```


MATRIX TILE SHAPE – MORE DETAILS

Program needs to multiply matrices: $C=A*B$

Example: C: 512x256, A: 512x1024, B: 1024x256 (#rows x #columns of elements)

- Ask multiplier for geometry information

$t1 \leftarrow (256 \ll 16) \mid 512$

vmts tc, t1, e32, mtc, rmo

- The multiplier replies with a tile shape stored in the scalar register *tc*.

For example, on a VLEN=512 system:


$tc = \{rows=16, cols=16, rmul=4, cmul=4, sew=32, mtype=mtc, cmo=0\}$

- If requested geometry is smaller than the supported size but can be processed by the multiplier (e.g. by filling in zeros): rows, cols shall correspond to the asked geometry.
- If the multiplier does not support the smaller geometry, return zero. → process in vector unit.

- We use *tc* to load a piece of matrix into the vector register group.

- The exact order of elements in the vector registers is not prescribed. It is left to the implementation.
- Portability can only be achieved with implementation specific tile load/store instructions.

```
// example
struct TileShape {
    unsigned int rows : 13;
    unsigned int cols : 13;
    unsigned int rmul  : 3;
    unsigned int cmul  : 3;
    unsigned int sew   : 3;
    unsigned int mtype : 2;
    unsigned int cmo   : 1;
};
```



WHAT ABOUT THE ACCUMULATOR?

- ✓ Case 1: Store directly to vector registers.
 - Simple and easy, more cycles and data movement.
 - OK for short vectors, small accumulators
- ✓ Case 2: treat accumulator like register bypass

✓ Register renaming

- For *vmmacc* *vd* is input & output register!
- No commit until next instruction uses register
- IF instruction is *vmmacc*
 - IF previous instruction was not *vmmacc*
 - Read *vd* into *Acc*
- ELSE IF previous instruction is *vmmacc*
 - Write *Acc* to *vd*

✓ “Reservation Station”

- Matrix multiply using *vd* lands in R.S.
 - Once *vd* is available: transfers data into *acc*
- When matrix multiply is finished: *vmmacc* remains in R.S.
 - Retired when subsequent *vmmacc* is issued, reusing the *vd*
 - If another instruction is issued which touches *vd*:
 - stores *acc* to *vd*
 - retires

RISC-V IME TG EVALUATION SCENARIOS

✓ A: Low-end “embedded” core

- $VLEN = 128$, $DLEN = 128$
- # of RF rd/wr ports = $3Rd + 1Wr$, possibly 2-way banked (plus Rd/Wr ports for St/Ld)
- # of cache rd/wr accesses = $1Rd/Wr$, VLEN-wide
- In-order, two-wide superscalar, no register renaming, ~2 GHz

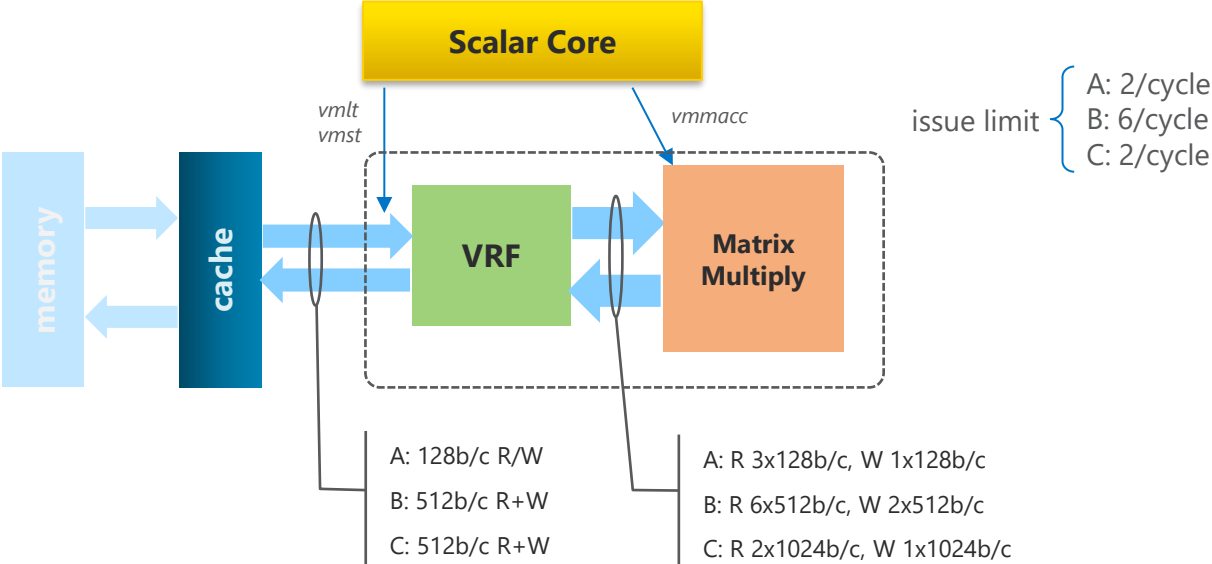
✓ B: High-end “apps” core

- $VLEN = 512$, $DLEN = 512$ (or 2×256)
- # of RF rd/wr ports = $6Rd + 2Wr$, possibly 2-way banked (plus Rd/Wr ports for St/Ld)
- # of cache rd/wr accesses = $1Rd + 1Wr$, VLEN-wide (or $2 \times VLEN/2$)
- Out-of-order, six-wide superscalar, full register renaming (speculative execution), ~3 GHz


✓ C: “HPC” core

- $VLEN = 16K$, $DLEN = 1K$ (as 16 64-bit-wide “lanes”)
- # of RF rd/wr ports = $32 \text{ 8B-wide Rd} + 16 \text{ 8B-wide Wr}$
- # of cache rd/wr accesses = $1Rd + 1Wr$, $DLEN/2=512b$ -wide
- Limited out-of-order, two-wide superscalar, limited register renaming (speculative execution), ~2 GHz

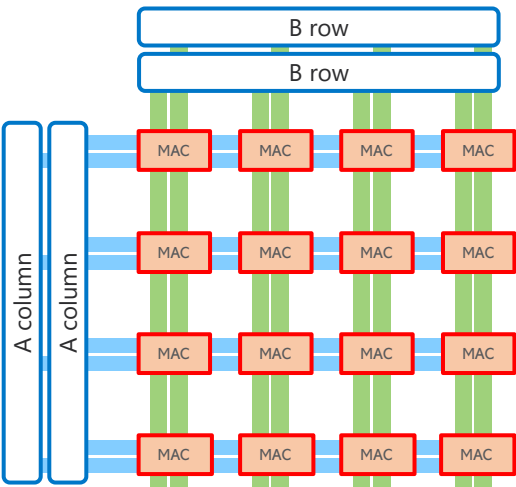
CONSTRAINTS




OUTER PRODUCT



Outer product
Rank-1 update





Outer product
Rank-2 update



Outer product
Accumulator larger than physical MAC array

INNER PRODUCT






$8 \cdot \text{DLEN} / \text{SEW}$ multiplies / cycle

Outer product with BW DLEN/cycle:
 $(\text{DLEN} / (2 \cdot \text{SEW}))^2$ multiplies / cycle
 - Can win if using full load BW for A,B
 - If $\text{DLEN} / \text{SEW} \geq 32$

OVERVIEW INNER VS. OUTER PRODUCT PERFORMANCE

1	Performance											
2	Core	Embedded			Apps			HPC			units	Note
3		i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64		
4												
5	outer	102	51	9	1024	512	128	10810	5405	676	GFLOPS/GOPS	
6	outer	170,7	85,3	12,8	1229	614	154	10923	5461	1008	GFLOPS/GOPS	register bypass approach
7		170,7	85,3	12,8	1229	614	154	10923	5461	1008	GFLOPS/GOPS	stride optimized VR loads
8												
9	outer 2	146,3	73,1	12,8	1229	614	192	10810	5405	676	GFLOPS/GOPS	
10	outer 2	170,7	85,3	21,3	1229	614	192	10923	5461	1008	GFLOPS/GOPS	register bypass approach
11	outer 2	256,0	128,0	21,3	2048	1024	256	10923	5461	1008	GFLOPS/GOPS	stride optimized VR loads
12												
13	inner	128	64	12,8	2048	1024	192	957	478	120	GFLOPS/GOPS	
14	inner	208,6	104,3	17,8	2973	1486	285	1982	991	248	GFLOPS/GOPS	cache Bs in dotproduct unit?
15												
16	geometry	8x8x8	8x4x8	4x2x8	16x16x16	16x8x16	8x4x16	64x128x128	64x64x128	64x16x64		
17	phys MAC	4x4	4x4	2x2	8x8	8x8	4x8	32x32	32x32	16x16		outer
18	phys MAC	4x8	4x8	2x4	8x16	8x16	8x8	32x32	32x32	16x16		outer 2

„HPC“ CORE



mtype	rows	cols	RMUL	CMUL
C	64	64	4	4
A	64	16	4	1
B	16	64	1	4

mtype	rows	cols	RMUL	CMUL
C	64	128	4	4
A	64	64	4	1
B	64	128	2	4




mtype	rows	cols	RMUL	CMUL
C	64	128	4	4
A	64	128	4	1
B	128	128	2	4

fp64 += fp64*fp64

fp32 += bf16*bf16

int32 += int8*int8

HIGH-END „APPS“ CORE



mtype	rows	cols	RMUL	CMUL
C	8	16	4	4
A	8	4	4	1
B	4	16	2	4

mtype	rows	cols	RMUL	CMUL
C	16	16	4	4
A	16	8	4	1
B	8	16	4	1




mtype	rows	cols	RMUL	CMUL
C	16	16	4	4
A	16	16	4	1
B	16	16	4	1

fp64 += fp64*fp64

fp32 += bf16*bf16

int32 += int8*int8

LOW-END „EMBEDDED“ CORE



mtype	rows	cols	RMUL	CMUL
C	4	8	4	4
A	4	2	4	1
B	2	8	2	4

fp64 += fp64*fp64

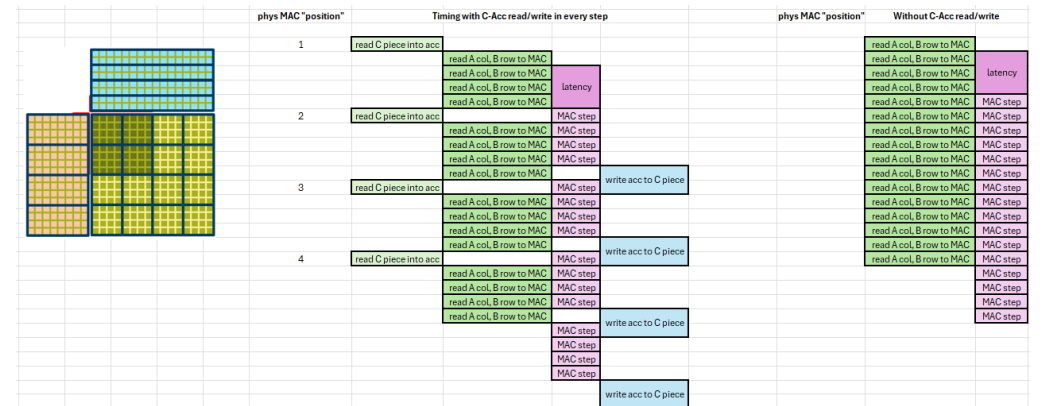
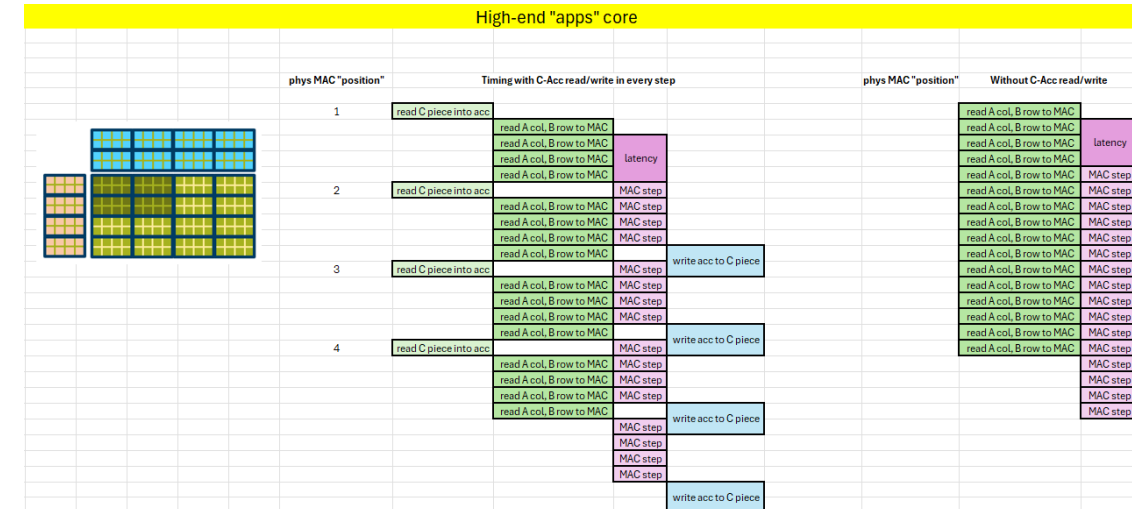
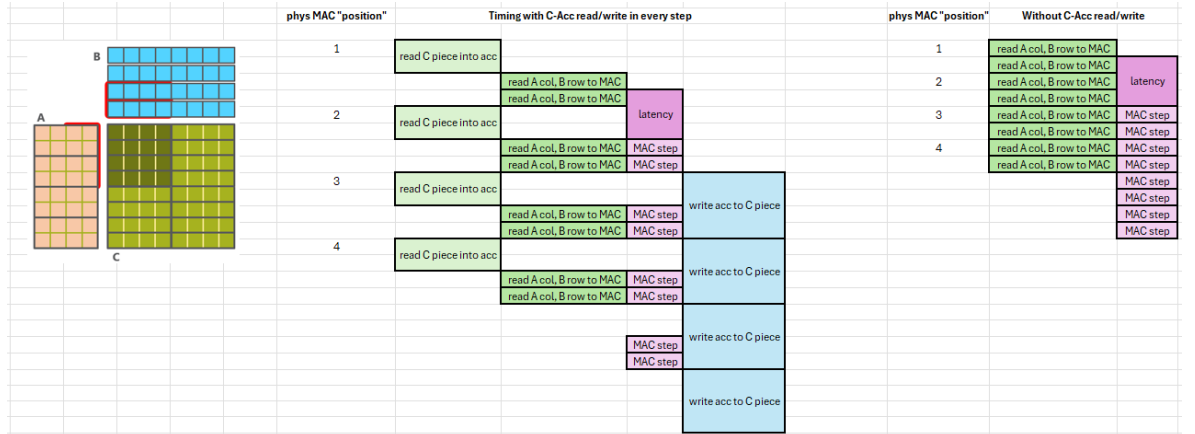
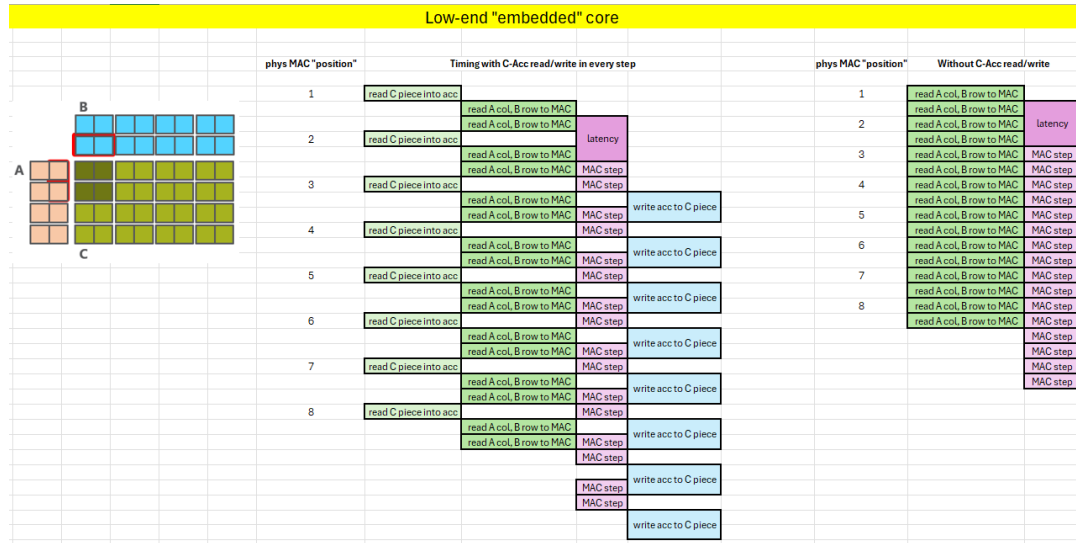
mtype	rows	cols	RMUL	CMUL
C	8	8	8	2
A	8	4	4	1
B	4	8	4	1

fp32 += bf16*bf16

mtype	rows	cols	RMUL	CMUL
C	8	8	8	2
A	4	2	4	1
B	2	8	2	4

int32 += int8*int8

EVALUATION - PIPELINES




EVALUATION – OUTER PRODUCT

1	Core	Embedded			Apps			HPC			units
2	VLEN	128	128	128	512	512	512	16384	16384	16384	bits
3	DLEN	128	128	128	512	512	512	1024	1024	1024	bits
4	Register read ports	3	3	3	6	6	6	2	2	2	ports of size DLEN
5	Register write ports	1	1	1	2	2	2	1	1	1	ports of size DLEN
6	Cache read bw	128	128	128	512	512	512	512	512	512	bits/cycle
7	Version	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	
8	MEW	32	32	64	32	32	64	32	32	64	bits (accumulator)
9	widening	4	2	1	4	2	1	4	2	1	factor
19	physical MAC array rows	4	4	2	8	8	4	32	32	16	MACs working in parallel
20	physical MAC array columns	4	4	2	8	8	8	32	32	16	MACs working in parallel
21	MAC latency	4	4	4	4	4	4	4	4	4	cycles
22											
28	Load all A,B VREGs from cache	12	12	12	20	20	16	384	384	256	cycles
31	read all phys MAC args A,B	2	2	2	4	4	4	32	32	16	cycles per position
32	acc size / phys MAC array sz	4	4	8	4	4	4	8	8	16	
33	Read all A,B args	8	8	16	16	16	16	256	256	256	cycles
34	exclude C read/write each step, register bypass approach										
35	one phys MAC array compute	2	2	2	4	4	4	32	32	16	cycles
36	all MAC arrays compute	12	12	20	20	20	20	260	260	260	cycles
37	number of MACs processed	512	256	64	4096	2048	512	1048576	524288	65536	
38	MACs/cycle	42,7	21,3	3,2	204,8	102,4	25,6	2730,7	1365,3	252,1	MACs/cycle
39	Mem BW for A,B load	11	11	10	26	26	38	64	64	63	Bytes/cycle
40	Frequency	2	2	2	3	3	3	2	2	2	GHz
41	Performance	171	85	13	1229	614	154	10923	5461	1008	GFLOPS/GOPS
42											
52	include C read/write each step, C write is being overlapped and hidden										
53	all MAC arrays compute	20	20	28	24	24	24	388	388	388	cycles
54	MACs/cycle	25,6	12,8	2,3	170,7	85,3	21,3	2702,5	1351,3	168,9	MACs/cycle
55	Mem BW for A,B load	6	6	7	21	21	32	63	63	42	Bytes/cycle
56	Performance	102	51	9	1024	512	128	10810	5405	676	GFLOPS/GOPS

EVALUATION – OUTER PRODUCT – 2 (INCREASED PHYS MAC ARRAY)

1	Core	Embedded			Apps			HPC			units
7	Version	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	
10	M	8	8	4	16	16	8	64	64	64	
11	K	8	4	2	16	8	4	128	64	16	
12	N	8	8	8	16	16	16	128	128	64	
13	RMUL_A	4	4	4	4	4	4	4	4	4	A row multiplier
14	CMUL_A	1	1	1	1	1	1	1	1	1	A column multiplier
15	RMUL_B	4	4	2	4	4	2	2	2	1	
16	CMUL_B	1	1	4	1	1	4	4	4	4	
17	RMUL_C	8	8	4	4	4	4	4	4	4	
18	CMUL_C	2	2	4	4	4	4	4	4	4	
19	physical MAC array rows	4	4	2	8	8	8	32	32	16	MACs working in parallel
20	physical MAC array columns	8	8	4	16	16	8	32	32	16	MACs working in parallel
21	MAC latency	4	4	4	4	4	4	4	4	4	cycles
28	Load all A,B VREGs from cache	12	12	12	20	20	16	384	384	256	cycles
29	Load all A,B (stride optimized)	8	8	12	8	8	12	384	384	256	cycles
32	read all phys MAC args A,B	2	2	2	4	4	4	32	32	16	cycles per position
33	acc size / phys MAC array sz	2	2	4	2	2	2	8	8	16	
34	Read all A,B args	4	4	8	8	8	8	256	256	256	cycles
35	exclude C read/write each step, register bypass approach										
36	one phys MAC array compute	2	2	2	4	4	4	32	32	16	cycles
37	all MAC arrays compute	8	8	12	12	12	12	260	260	260	cycles
38	number of MACs processed	512	256	64	4096	2048	512	1048576	524288	65536	
39	MACs/cycle	42,7	21,3	5,3	204,8	102,4	32,0	2730,7	1365,3	252,1	MACs/cycle
40	Mem BW for A,B load	11	11	16	26	26	48	64	64	63	Bytes/cycle
41	Frequency	2	2	2	3	3	3	2	2	2	GHz
42	Performance	171	85	21	1229	614	192	10923	5461	1008	GFLOPS/GOPS
43											
53	include C read/write each step, C write is being overlapped and hidden										
54	all MAC arrays compute	14	14	20	20	20	16	388	388	388	cycles
55	MACs/cycle	36,6	18,3	3,2	204,8	102,4	32,0	2702,5	1351,3	168,9	MACs/cycle
56	Mem BW for A,B load	9	9	10	26	26	48	63	63	42	Bytes/cycle
57	Performance	146	73	13	1229	614	192	10810	5405	676	GFLOPS/GOPS
58											
59	exclude C read/write each step, register bypass approach, stride optimized VR loads										
60	MACs/cycle	64,0	32,0	5,3	341,3	170,7	42,7	2730,7	1365,3	252,1	MACs/cycle
61	Mem BW for A,B load	16	16	16	43	43	64	64	64	63	Bytes/cycle
62	Performance	256	128	21	2048	1024	256	10923	5461	1008	GFLOPS/GOPS

INNER PRODUCT: LOW-END „EMBEDDED“ CORE



mtype	rows	cols	RMUL	CMUL
C	1	8	1	4
A	1	2	1	1
B	2	8	1	8

fp64 += fp64*fp64

mtype	rows	cols	RMUL	CMUL
C	1	8	1	2
A	1	8	1	1
B	8	8	1	8

fp32 += bf16*bf16

mtype	rows	cols	RMUL	CMUL
C	1	8	1	2
A	1	16	1	1
B	16	8	1	8

int32 += int8*int8

EVALUATION – INNER PRODUCT

1	Core	Embedded			Apps			HPC			units
2	VLEN	128	128	128	512	512	512	16384	16384	16384	bits
3	DLEN	128	128	128	512	512	512	1024	1024	1024	bits
4	Register read ports	3	3	3	6	6	6	2	2	2	ports of size DLEN
5	Register write ports	1	1	1	2	2	2	1	1	1	ports of size DLEN
6	Cache read bw	128	128	128	512	512	512	512	512	512	bits/cycle
7	Issues perc cycle	2	2	2	6	6	6	2	2	2	max instr issues
8	Version	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	
9	MEW	32	32	64	32	32	64	32	32	64	bits (accumulator)
10	widening	4	2	1	4	2	1	4	2	1	factor
11											
12	cols of B	8	8	8	16	16	8	16	16	16	
13	rows of A	11	11	5	15	15	23	15	15	15	"iterations"
14	width of C (elements)	8	8	8	16	16	8	16	16	16	
15	width of C (VRs)	2	2	4	1	1	1	0,03125	0,03125	0,0625	# of VRs
16	A, B elements per VL	16	8	2	64	32	8	2048	1024	256	
17	MACCs	1408	704	80	15360	7680	1472	491520	245760	61440	multiply-accumulate
18	width of dotprod	128	128	128	512	512	512	1024	1024	1024	bits
19											
20	Instructions	54	54	36	97	97	90	97	97	97	instructions
21	Min issue time	27	27	18	17	17	15	49	49	49	cycles
22	Loads (vle)	19	19	13	31	31	31	496	496	496	DLEN bits
23	Load time	19	19	13	31	31	31	992	992	992	cycles
24	VRF reads per vmacc instr	1408	1408	1664	9216	9216	5120	279040	279040	279552	bits
25	time for1 instr VRF reads	4	4	5	3	3	2	137	137	137	cycles
26	Cycles all VRF reads	44	44	25	45	45	46	2055	2055	2055	cycles
27	Time one iter dotprod	1	1	1	1	1	1	16	16	16	cycles
28	Time all iter dotproducts	11	11	5	15	15	23	240	240	240	cycles
29											
30	Max(issue,load,VRFrd,dotp)	44	44	25	45	45	46	2055	2055	2055	cycles
31	Ops/cycle	64,0	32,0	6,4	682,7	341,3	64,0	478,4	239,2	59,8	
32	Ops/cycle no VRF	104,3	52,1	8,9	991,0	495,5	95,0	991,0	495,5	123,9	
33	Ops/load	148,2	74,1	12,3	991,0	495,5	95,0	991,0	495,5	123,9	
34	Frequency	2	2	2	3	3	3	2	2	2	GHz
35	Performance	128,0	64,0	12,8	2048,0	1024,0	192,0	956,7	478,4	119,6	GFLOPS/GOPS
36	Performance no VRF	208,6	104,3	17,8	2972,9	1486,5	284,9	1981,9	991,0	247,7	GFLOPS/GOPS

Match ~ Option F
evaluation



EVALUATION – INNER PRODUCT

1	Core	Embedded			Apps			HPC			units
2	VLEN	128	128	128	512	512	512	16384	16384	16384	bits
3	DLEN	128	128	128	512	512	512	1024	1024	1024	bits
4	Register read ports	3	3	3	6	6	6	2	2	2	ports of size DLEN
5	Register write ports	1	1	1	2	2	2	1	1	1	ports of size DLEN
6	Cache read bw	128	128	128	512	512	512	512	512	512	bits/cycle
7	Issues perc cycle	2	2	2	6	6	6	2	2	2	max instr issues
8	Version	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	i8*i8+=i32	b16*b16+=f32	f64*f64+=f64	
9	MEW	32	32	64	32	32	64	32	32	64	bits (accumulator)
10	widening	4	2	1	4	2	1	4	2	1	factor
11											
12	cols of B	8	8	8	16	16	8	16	16	16	
13	rows of A	11	11	5	15	15	23	15	15	15	"iterations"
14	width of C (elements)	8	8	8	16	16	8	16	16	16	
15	width of C (VRs)	2	2	4	1	1	1	0,03125	0,03125	0,0625	# of VRs
16	A, B elements per VL	16	8	2	64	32	8	2048	1024	256	
17	MACCs	1408	704	80	15360	7680	1472	491520	245760	61440	multiply-accumulate
18	width of dotprod	128	128	128	512	512	512	1024	1024	1024	bits
19											
20	Instructions	54	54	36	97	97	90	97	97	97	instructions
1	Min issue time	27	27	18	17	17	15	49	49	49	cycles
2	Loads (vle)	19	19	13	31	31	31	496	496	496	DLEN bits
3	Load time	19	19	13	31	31	31	992	992	992	cycles
4	VRF reads per vmacc instr	1408	1408	1664	9216	9216	5120	279040	279040	279552	bits
5	time for 1 instr VRF reads	4	4	5	3	3	2	137	137	137	cycles
6	Cycles all VRF reads	44	44	25	45	45	46	2055	2055	2055	cycles
7	Time one iter dotprod	1	1	1	1	1	1	16	16	16	cycles
28	Time all iter dotproducts	11	11	5	15	15	23	240	240	240	cycles
29											
30	Max(issue,load,VRFrd,dotp)	44	44	25	45	45	46	2055	2055	2055	cycles
31	Ops/cycle	64,0	32,0	6,4	682,7	341,3	64,0	478,4	239,2	59,8	
32	Ops/cycle no VRF	104,3	52,1	8,9	991,0	495,5	95,0	991,0	495,5	123,9	
33	Ops/load	148,2	74,1	12,3	991,0	495,5	95,0	991,0	495,5	123,9	
34	Frequency	2	2	2	3	3	3	2	2	2	GHz
35	Performance	128,0	64,0	12,8	2048,0	1024,0	192,0	956,7	478,4	119,6	GFLOPS/GOPS
36	Performance no VRF	208,6	104,3	17,8	2972,9	1486,5	284,9	1981,9	991,0	247,7	GFLOPS/GOPS

Clarified by Jose during the meeting: „F“ uses banking of VRF (2x256b).

Implies that B VRs are not transferred into the fat dotprod unit each time they are used.

Match ~ Option F evaluation

$$C = \text{ALPHA} \cdot A \cdot B + \text{BETA} \cdot C$$

Matrix data types: A,B,C: FP32

```

vmts X(tc),X(M|N<<16),e32,mtc,rmo
vmts X(ta),X(M|K<<16),e32,mta,rmo
vmts X(tb),X(K|N<<16),e32,mtb,rmo
IF (RMUL(tc)*CMUL(tc)==1 &&
    RMUL(ta)*CMUL(ta)==1 &&
    RMUL(tb)*CMUL(tb)==1) THEN          # 16 steps
    call gemm_16_steps_1_1_1
ELSE IF (RMUL(tc)*CMUL(tc)==16) THEN    # one step
    IF (RMUL(ta)*CMUL(ta)==4) THEN
        IF (RMUL(tb)*CMUL(tb)==4) THEN
            call gemm_one_step_16_4_4
        ELSE IF (RMUL(tb)*CMUL(tb)==8) THEN
            call gemm_one_step_16_4_8
        ELSE
            ...
        ENDIF
    ENDIF
ENDIF

```

gemm_one_step_16_4_4:

```

tm = ROWS(tc); tn = COLS(tc); tk = ROWS(tb)
for m = 0, M, m += tm do
    for n = 0, N, n += tn do
        vsetvli zero,x(-1),e32,m8,ta,ma
        vfmv.v.f v16,f(0.0f)
        vfmv.v.f v24,f(0.0f)
        vmmset(tc,ta,tb)
        for k = 0, K, k += tk do
            vmlt v8,x(&A[m,k]),x(rstr_A|cstr_A<<31|mta<<62)
            vmlt v12,x(&B[k,n]),x(rstr_B|cstr_B<<31|mtb<<62)
            vfmmacc v16,v8,v12
        vmlt v0,x(&C[m,n]),x(rstr_C|cstr_C<<31|mtc<<62)
        vsetvli zero,X(-1),e32,m8,ta,ma
        vfmul.vf v16,v16,alpha
        vfmul.vf v24,v24,alpha
        vfmaccc.vf v16,v0,beta
        vfmaccc.vf v24,v8,beta
        vmst v16,x(&C[m,n]),x(rstr_C|cstr_C<<31|mtc<<62)
    endfor
endfor

```

SUMMARY

- ✓ ISA can be portable
 - **Need flexible query of hardware capabilities**
 - **Tile load/store commands**
- ✓ Implementor has the choice
 - Choose hardware (sequencer) vs. software balance
 - Inner or outer product
 - Optimizations: buffering B or Accumulator or ...
- ✓ Can map most of the approaches
 - Segments of A, C tiles not yet handled
 - (Software should know if approach is weight stationary or output stationary)

